

# logistic\_regularization

May 29, 2026

```
[13]: %load_ext autoreload
      %autoreload 2
```

The autoreload extension is already loaded. To reload it, use:  
%reload\_ext autoreload

```
[70]: import numpy as np
      import matplotlib.pyplot as plt

      from sklearn.datasets import make_moons, make_circles
      from scipy.optimize import minimize
      from scipy.special import expit, logit

      import warnings
      warnings.filterwarnings("ignore")

      import self_training as st
```

## 1 Analysis of Self-Training in Simple Logistic Regression

```
[96]: cfg = st.SelfTrainingConfig(
      N=100, M=10000, N_test=10000, d=500,
      mu_plus=0.1*np.ones(500), mu_minus=-0.1*np.ones(500),
      p=0.9, T=10, n_runs=20, n_jobs=10, seed=0,
      sigma_scale=2.0
      )

      rng = np.random.default_rng(cfg.seed)
      X_lab, Y_lab = st.sample_data(cfg.N, cfg, rng)
      X_unl, _ = st.sample_data(cfg.M, cfg, rng)
      X_test, Y_test = st.sample_data(cfg.N_test, cfg, rng)

      T = 50
```

```
[ ]: def res_logistic_regression_N(cfg, runs, seeds, penalty, C, total, T, p, N_list):
      length = len(N_list)
```

```

res = np.zeros((length, runs, T+1))

for i in range(length):
    for j in range(runs):

        seed = seeds[j]
        rng = np.random.default_rng(seed)
        X, Y = st.sample_data(total, cfg, rng)
        X_test, Y_test = st.sample_data(cfg.N_test, cfg, rng)

        N = N_list[i]
        X_lab, Y_lab = X[:N], Y[:N]
        X_unl = X[N:]

        res[i, j] = st.run_self_training(p, T, X_lab, Y_lab, X_unl, X_test,
↪Y_test, penalty=penalty, C=C)

    return res

def plot_logistic_regression_N(res, runs, penalty, C, total, T, p, N_list):
    length = len(N_list)
    cmap = plt.colormaps['plasma']
    fig, axes = plt.subplots(1, length, figsize=(15, 4), sharey=True)
    for i in range(length):
        N = N_list[i]

        results = res[i]

        mean_trajectory = np.mean(results, axis=0)
        median_trajectory = np.median(results, axis=0)
        lower_bound = np.percentile(results, 10, axis=0) # 10th percentile
        upper_bound = np.percentile(results, 90, axis=0)

        color_index = i / length
        line_color = cmap(color_index)

        axes[i].plot(np.arange(T+1), median_trajectory, color=line_color,
                     linewidth=1.0, label=f"N={N},  $\lambda = {(total-N)/N:.2f}")

        axes[i].plot(np.arange(T+1), mean_trajectory, color=line_color,
                     linewidth=1.0, linestyle='--', alpha=0.8)

        axes[i].fill_between(np.arange(T+1),
                             lower_bound,
                             upper_bound,
                             color=line_color, alpha=0.2)$ 
```

```

        axes[i].grid()
        axes[i].legend()
        axes[i].set_xlabel(r"Iterations")

    axes[0].set_ylabel(r"Error")
    fig.suptitle(f"Test Error of Self-Training with $N+M={total}$, $p={p}$, $
    ↪penalty={penalty}$, $C={C}$, averaged over {runs} runs")
    plt.show()

```

```

[119]: def res_logistic_regression_p(cfg, runs, seeds, penalty, C, total, T, N,
    ↪p_list):
    length = len(p_list)
    res = np.zeros((length, runs, T+1))

    for i in range(length):
        for j in range(runs):

            seed = seeds[j]
            rng = np.random.default_rng(seed)
            X, Y = st.sample_data(total, cfg, rng)
            X_test, Y_test = st.sample_data(cfg.N_test, cfg, rng)

            p = p_list[i]
            X_lab, Y_lab = X[:N], Y[:N]
            X_unl = X[N:]

            res[i, j] = st.run_self_training(p, T, X_lab, Y_lab, X_unl, X_test,
    ↪Y_test, penalty=penalty, C=C)

    return res

def plot_logistic_regression_p(res, runs, penalty, C, total, T, N, p_list):
    length = len(p_list)
    cmap = plt.colormaps['plasma']
    fig, axes = plt.subplots(1, length, figsize=(15, 4), sharey=True)
    for i in range(length):
        p = p_list[i]

        results = res[i]

        mean_trajectory = np.mean(results, axis=0)
        median_trajectory = np.median(results, axis=0)
        lower_bound = np.percentile(results, 10, axis=0) # 10th percentile
        upper_bound = np.percentile(results, 90, axis=0)

        color_index = i / length
        line_color = cmap(color_index)

```

```

axes[i].plot(np.arange(T+1), median_trajectory, color=line_color,
             linewidth=1.0, label=f"p={p}")

axes[i].plot(np.arange(T+1), mean_trajectory, color=line_color,
             linewidth=1.0, linestyle='--', alpha=0.8)

axes[i].fill_between(np.arange(T+1),
                    lower_bound,
                    upper_bound,
                    color=line_color, alpha=0.2)

axes[i].grid()
axes[i].legend()
axes[i].set_xlabel(r"Iterations")

axes[0].set_ylabel(r"Error")
fig.suptitle(f"Test Error of Self-Training with $N={total}$, $M={total-N}$,  $\hookrightarrow$ 
 $\hookrightarrow$  $p={p}$, penalty={penalty}, $C={C}$, averaged over {runs} runs")
plt.show()

```

```

[120]: runs = 10
       seeds = range(runs)
       total = 10000

```

## 1.1 No Regularization

```

[121]: penalty = None
       C = 1.0



       # Varying N
       p = 0.99
       N_list = [100, 500, 1000, 5000, 8000]

       N_res = res_logistic_regression_N(cfg, runs, seeds, penalty, C, total, T, p,  $\hookrightarrow$ 
 $\hookrightarrow$  N_list)
       plot_logistic_regression_N(N_res, runs, penalty, C, total, T, p, N_list)

       # Varying p
       N = 100
       p_list = [0.5, 0.75, 0.9, 0.95, 0.99]

       p_res = res_logistic_regression_p(cfg, runs, seeds, penalty, C, total, T, N,  $\hookrightarrow$ 
 $\hookrightarrow$  p_list)
       plot_logistic_regression_p(p_res, runs, penalty, C, total, T, N, p_list)

```



## 1.2 $L^2$ Regularization

$$C = 1/\lambda$$

$$C\ell(YX^\top\beta) + \frac{1}{2}\|\beta\|_2^2$$

```
[126]: penalty = 'l2'
C = 0.5



# Varying N
p = 0.99
N_list = [100, 500, 1000, 5000, 8000]

N_res = res_logistic_regression_N(cfg, runs, seeds, penalty, C, total, T, p, N_list)
plot_logistic_regression_N(N_res, runs, penalty, C, total, T, p, N_list)

# Varying p
N = 100
```

```
p_list = [0.5, 0.75, 0.9, 0.95, 0.99]
```

```
p_res = res_logistic_regression_p(cfg, runs, seeds, penalty, C, total, T, N,
    ↪ p_list)
plot_logistic_regression_p(p_res, runs, penalty, C, total, T, N, p_list)
```



```
[124]: penalty = 'l2'
C = 0.1

# Varying N
p = 0.99
N_list = [100, 500, 1000, 5000, 8000]



N_res = res_logistic_regression_N(cfg, runs, seeds, penalty, C, total, T, p,
    ↪ N_list)
plot_logistic_regression_N(N_res, runs, penalty, C, total, T, p, N_list)

# Varying p
N = 100
p_list = [0.5, 0.75, 0.9, 0.95, 0.99]
```

```

p_res = res_logistic_regression_p(cfg, runs, seeds, penalty, C, total, T, N,
    ↪ p_list)
plot_logistic_regression_p(p_res, runs, penalty, C, total, T, N, p_list)

```



```

[125]: penalty = 'l2'
C = 0.02

# Varying N
p = 0.99
N_list = [100, 500, 1000, 5000, 8000]

N_res = res_logistic_regression_N(cfg, runs, seeds, penalty, C, total, T, p,
    ↪ N_list)
plot_logistic_regression_N(N_res, runs, penalty, C, total, T, p, N_list)



# Varying p
N = 100
p_list = [0.5, 0.75, 0.9, 0.95, 0.99]

```

```

p_res = res_logistic_regression_p(cfg, runs, seeds, penalty, C, total, T, N,
    ↪ p_list)
plot_logistic_regression_p(p_res, runs, penalty, C, total, T, N, p_list)

```



## 2 Analysis of Self-Training in Series Logistic Regression and Series Density Estimation (B-Splines)

```

[20]: class PrecomputedAdditiveLegendre:
    def __init__(self, X_global: np.ndarray, K_max: int, quad_deg: int = 200):
        """
        Precomputes the complete basis design matrices for the entire dataset
        and the quadrature nodes required for numerical integration.
        """
        self.K_max = K_max
        self.d = X_global.shape[1] if X_global.ndim > 1 else 1

        # 1. Determine scaling bounds to map all features into [-1, 1]
        self.x_min = np.min(X_global, axis=0)
        self.x_max = np.max(X_global, axis=0)

```

```

self.scale = self.x_max - self.x_min
self.scale[self.scale == 0] = 1.0 # Prevent division by zero

# 2. Quadrature setup for integration over [-1, 1]
nodes_1d, weights_1d = np.polynomial.legendre.leggauss(quad_deg)
self.nodes_1d = nodes_1d
self.weights_1d = weights_1d

# 3. Precompute Global Data Matrix \Psi_global
self.Psi_global = self._build_additive_psi(X_global)

# 4. Precompute Nodes Matrix \Psi_nodes & \Lambda
# Since we use an additive model over d dimensions, the integral splits:
# \int \sum L_k(x^{(j)}) dx = d \int L_k(x) dx (assuming independent,
→uniform measure)
self.Psi_nodes = self._build_1d_psi(self.nodes_1d) * self.d
self.Lambda_global = self.Psi_nodes.T @ self.weights_1d

def _build_1d_psi(self, x_1d: np.ndarray) -> np.ndarray:
    """Helper: Builds 1D Legendre basis."""
    Psi = np.zeros((x_1d.shape[0], self.K_max))
    for k in range(self.K_max):
        coef = np.zeros(k + 1)
        coef[k] = 1.0
        Psi[:, k] = np.polynomial.legendre.legval(x_1d, coef)
    return Psi

def _build_additive_psi(self, X: np.ndarray) -> np.ndarray:
    """Helper: Builds Additive Legendre basis across d dimensions."""
    X_scaled = 2.0 * (X - self.x_min) / self.scale - 1.0
    if self.d == 1:
        return self._build_1d_psi(X_scaled.squeeze())

    N = X.shape[0]
    Psi = np.zeros((N, self.K_max))
    for k in range(self.K_max):
        coef = np.zeros(k + 1)
        coef[k] = 1.0
        # Evaluate polynomial on each dimension and sum (Additive Model)
        evaluated = np.polynomial.legendre.legval(X_scaled, coef)
        Psi[:, k] = np.sum(evaluated, axis=1)
    return Psi

# Series Logistic Regression

def fit_logistic_offset(Psi: np.ndarray, Y: np.ndarray, offset: np.ndarray,

```

```

        l2_reg: float = 1e-4) -> np.ndarray:
    """Solves for \vartheta minimizing the empirical logistic loss with an
    ↪offset."""
    N = Psi.shape[0]

    def objective(vartheta):
        logits = Psi @ vartheta + offset
        loss = np.mean(np.logaddexp(0, -Y * logits))
        return loss + 0.5 * l2_reg * np.sum(vartheta**2)

    def gradient(vartheta):
        logits = Psi @ vartheta + offset
        grad = -Psi.T @ (Y * (1 - expit(Y * logits))) / N
        return grad + l2_reg * vartheta

    res = minimize(objective, np.zeros(Psi.shape[1]), jac=gradient,
    ↪method='L-BFGS-B')
    return res.x

def run_series_logistic_regression(X_L0: np.ndarray, Y_L0: np.ndarray,
                                  X_U0: np.ndarray, X_test: np.ndarray, Y_test:
    ↪ np.ndarray,
                                  kappa: float, T: int, K_seq: list, l2_reg:
    ↪float = 1e-4) -> np.ndarray:
    """
    Returns an array of shape (T+1,) containing the test error at each round.
    """
    N, M = X_L0.shape[0], X_U0.shape[0]
    K_max = K_seq[-1]

    # 1. Global Precomputation
    X_global = np.vstack([X_L0, X_U0])
    basis = PrecomputedAdditiveLegendre(X_global, K_max)
    Psi_global = basis.Psi_global

    # Precompute Test Matrix (using scaling bounds from training data)
    Psi_test = basis._build_additive_psi(X_test)

    idx_L0 = np.arange(N)
    idx_U0 = np.arange(N, N + M)

    # Initialization t = 0
    K_0 = K_seq[0]
    Psi_L0_J0 = Psi_global[idx_L0, :K_0]
    theta_current = fit_logistic_offset(Psi_L0_J0, Y_L0, np.zeros(N), l2_reg)

    # Track Error at t=0

```

```

Y_pred_0 = np.sign(Psi_test[:, :K_0] @ theta_current)
Y_pred_0[Y_pred_0 == 0] = 1 # Handle exact 0 logits
errors = [np.mean(Y_pred_0 != Y_test)]

for t in range(T):
    K_prev = K_seq[t]
    K_next = K_seq[t+1]

    # 1. Evaluate current function on U0
    h_U0 = Psi_global[idx_U0, :K_prev] @ theta_current

    # 2. Pseudo-labeling
    mask = np.abs(h_U0) >= kappa

    if mask.sum() > 0:
        idx_pseudo = idx_U0[mask]
        Y_pseudo = np.sign(h_U0[mask])
        # Construct  $\mathcal{L}_{t+1}$  (LO U Pseudo)
        idx_Lt1 = np.concatenate([idx_L0, idx_pseudo])
        Y_Lt1 = np.concatenate([Y_L0, Y_pseudo])
    else:
        # If no new points, still fit next basis block on L0 to reduce
        → approximation error
        idx_Lt1 = np.copy(idx_L0)
        Y_Lt1 = np.copy(Y_L0)

    # 3. Slicing for  $I_t$  (the new basis columns)
    Psi_Lt1_I = Psi_global[idx_Lt1, K_prev:K_next]

    # 4. Compute Offset  $h_{t-1}(X)$ 
    offset_Lt1 = Psi_global[idx_Lt1, :K_prev] @ theta_current

    # 5. Fit new block  $\vartheta_{t+1}$ 
    vartheta = fit_logistic_offset(Psi_Lt1_I, Y_Lt1, offset_Lt1, l2_reg)

    # 6. Concatenate  $\theta_{t+1}$ 
    theta_current = np.concatenate([theta_current, vartheta])

    # 7. Evaluate Error for round t+1
    Y_pred_t1 = np.sign(Psi_test[:, :K_next] @ theta_current)
    Y_pred_t1[Y_pred_t1 == 0] = 1
    errors.append(np.mean(Y_pred_t1 != Y_test))

return np.array(errors)

# Series Log-density Estimation

```

```

def solve_density_update(Psi_data: np.ndarray, Psi_nodes: np.ndarray,
                        base_h_nodes: np.ndarray, Lambda_I: np.ndarray,
                        weights: np.ndarray) -> np.ndarray:
    """
    Implements the Solve() operator with analytical Jacobians.
    """
    N = Psi_data.shape[0]
    if N == 0:
        return np.zeros(Psi_data.shape[1])

    def objective(vartheta):
        h_nodes = base_h_nodes + Psi_nodes @ vartheta
        max_h = np.max(h_nodes)

        # log \int \exp(h(y)) dy via log-sum-exp
        E = np.exp(h_nodes - max_h)
        Z = np.sum(weights * E)
        psi_J = max_h + np.log(Z)

        empirical_mean = np.mean(Psi_data @ vartheta)
        return -(empirical_mean - psi_J)

    def gradient(vartheta):
        h_nodes = base_h_nodes + Psi_nodes @ vartheta
        max_h = np.max(h_nodes)

        E = np.exp(h_nodes - max_h)
        Z = np.sum(weights * E)

        # Analytical gradient
        expected_Psi = (Psi_nodes.T @ (weights * E)) / Z
        empirical_Psi = np.mean(Psi_data, axis=0)
        return -(empirical_Psi - expected_Psi)

    # Strict linear constraint: \Lambda_I^top \vartheta = 0
    constraints = [{
        'type': 'eq',
        'fun': lambda v: Lambda_I @ v,
        'jac': lambda v: Lambda_I
    }]

    res = minimize(objective, np.zeros(Psi_data.shape[1]), jac=gradient,
                  method='SLSQP', constraints=constraints, options={'maxiter': 500})
    return res.x

```

```

def run_series_log_density(X_L0: np.ndarray, Y_L0: np.ndarray,
                           X_U0: np.ndarray, X_test: np.ndarray, Y_test: np.
                           ↪ ndarray,
                           pi: float, kappa: float, T: int, K_seq: list) -> np.
                           ↪ ndarray:
    """
    Returns an array of shape (T+1,) containing the test error at each round.
    """
    N, M = X_L0.shape[0], X_U0.shape[0]
    K_max = K_seq[-1]

    # 1. Global Precomputation
    X_global = np.vstack([X_L0, X_U0])
    basis = PrecomputedAdditiveLegendre(X_global, K_max)

    Psi_global = basis.Psi_global
    Psi_nodes = basis.Psi_nodes
    weights = basis.weights_1d
    Lambda_global = basis.Lambda_global

    Psi_test = basis._build_additive_psi(X_test)

    idx_L0_plus = np.where(Y_L0 == 1)[0]
    idx_L0_minus = np.where(Y_L0 == -1)[0]
    idx_U0 = np.arange(N, N + M)

    # Initialization t = 0
    K_0 = K_seq[0]
    theta_plus = solve_density_update(
        Psi_global[idx_L0_plus, :K_0], Psi_nodes[:, :K_0], np.
        ↪ zeros(len(weights)),
        Lambda_global[:K_0], weights)

    theta_minus = solve_density_update(
        Psi_global[idx_L0_minus, :K_0], Psi_nodes[:, :K_0], np.
        ↪ zeros(len(weights)),
        Lambda_global[:K_0], weights)

    def evaluate_test_error(t_plus, t_minus, K_current):
        P_nodes = Psi_nodes[:, :K_current]
        P_test = Psi_test[:, :K_current]

        Z_plus = np.sum(weights * np.exp(P_nodes @ t_plus))
        Z_minus = np.sum(weights * np.exp(P_nodes @ t_minus))

```

```

    log_odds = P_test @ (t_plus - t_minus) - np.log(Z_plus) + np.
↪ log(Z_minus) + np.log(pi / (1 - pi))
    preds = np.sign(log_odds)
    preds[preds == 0] = 1
    return np.mean(preds != Y_test)

# Track Error at t=0
errors = [evaluate_test_error(theta_plus, theta_minus, K_0)]

def classifier_logits(idx_data, K_current, t_plus, t_minus):
    P_data = Psi_global[idx_data, :K_current]
    P_nodes = Psi_nodes[:, :K_current]

    Z_plus = np.sum(weights * np.exp(P_nodes @ t_plus))
    Z_minus = np.sum(weights * np.exp(P_nodes @ t_minus))

    log_odds = P_data @ (t_plus - t_minus) - np.log(Z_plus) + np.
↪ log(Z_minus) + np.log(pi / (1 - pi))
    return log_odds

for t in range(T):
    K_prev = K_seq[t]
    K_next = K_seq[t+1]

    # 1. Pseudo-labeling
    c_U0 = classifier_logits(idx_U0, K_prev, theta_plus, theta_minus)
    mask = np.abs(c_U0) >= kappa

    if mask.sum() > 0:
        idx_pseudo_plus = idx_U0[mask][np.sign(c_U0[mask]) == 1]
        idx_pseudo_minus = idx_U0[mask][np.sign(c_U0[mask]) == -1]
        # Construct  $\mathcal{L}_{t+1}^{pm}$ 
        idx_Lt1_plus = np.concatenate([idx_L0_plus, idx_pseudo_plus])
        idx_Lt1_minus = np.concatenate([idx_L0_minus, idx_pseudo_minus])
    else:
        # Fallback to unaugmented dataset to expand basis
        idx_Lt1_plus = np.copy(idx_L0_plus)
        idx_Lt1_minus = np.copy(idx_L0_minus)

    # 3. Base functions for nodes
    base_h_nodes_plus = Psi_nodes[:, :K_prev] @ theta_plus
    base_h_nodes_minus = Psi_nodes[:, :K_prev] @ theta_minus

    # 4. Slicing for new block  $I_t$ 
    Psi_I_nodes = Psi_nodes[:, K_prev:K_next]
    Lambda_I = Lambda_global[K_prev:K_next]

```

```

# 5. Fit  $\vartheta_{t+1}^{\pm}$ 
vartheta_plus = solve_density_update(
    Psi_global[idx_Lt1_plus, K_prev:K_next], Psi_I_nodes,
    base_h_nodes_plus, Lambda_I, weights)

vartheta_minus = solve_density_update(
    Psi_global[idx_Lt1_minus, K_prev:K_next], Psi_I_nodes,
    base_h_nodes_minus, Lambda_I, weights)

# 6. Update parameters
theta_plus = np.concatenate([theta_plus, vartheta_plus])
theta_minus = np.concatenate([theta_minus, vartheta_minus])

# 7. Evaluate Error for round t+1
errors.append(evaluate_test_error(theta_plus, theta_minus, K_next))

return np.array(errors)

```

```

[91]: def results(X_L0, Y_L0, X_U0, X_test, Y_test, pi, p, T, K_seq):
    kappa = logit(p)
    d = X_L0.shape[1]
    print(f"Parameters: d={d}, T={T}, p={p}")
    print(f"\n--- Series Logistic Regression ---")
    series_logistic_error = run_series_logistic_regression(
        X_L0=X_L0, Y_L0=Y_L0, X_U0=X_U0, X_test=X_test, Y_test=Y_test,
    ↪ kappa=kappa, T=T, K_seq=K_seq, l2_reg=1e-3
    )
    print(series_logistic_error)

    print(f"\n--- Series Log-density Estimation ---")
    series_density_error = run_series_log_density(
        X_L0=X_L0, Y_L0=Y_L0, X_U0=X_U0, X_test=X_test, Y_test=Y_test, pi=pi,
    ↪ kappa=kappa, T=T, K_seq=K_seq
    )
    print(series_density_error)

    print(f"\n--- Simple Logistic Regression ---")
    simple_logistic_error = st.run_self_training(
        p, T, X_lab=X_L0, Y_lab=Y_L0, X_unl=X_U0, X_test=X_test, Y_test=Y_test
    )
    print(simple_logistic_error)

```

```

[92]: N = 100
total = 10000
M = total - N
N_test = 5000
pi = 0.5

```

```

T = 3
seed = 0
sigma_scale = 2.0

# Sieve Expansion Schedule: Define the number of basis functions K at each
↳round t.
# For T=3 rounds, we need 4 values: [K_0, K_1, K_2, K_3].
# Here, we start with 3 basis functions and expand by 2 in each self-training
↳round.
K_seq = [3, 5, 7, 9]

```

## 2.1 Gaussian Mixture in High Dimensional Setting (d=500)

```

[93]: d = 500
p = 0.99
cfg = st.SelfTrainingConfig(
    N=N, M=M, N_test=N_test, d=d,
    mu_plus=0.1 * np.ones(d), mu_minus=-0.1 * np.ones(d),
    p=p, T=T, n_runs=1, n_jobs=1, seed=seed,
    sigma_scale=sigma_scale
)

rng = np.random.default_rng(seed)
X, Y = st.sample_data(total, cfg, rng)
X_test, Y_test = st.sample_data(N_test, cfg, rng)
X_L0 = X[:N]
Y_L0 = Y[:N]
X_U0 = X[N:total]

results(X_L0, Y_L0, X_U0, X_test, Y_test, pi, p, T, K_seq)

```

Parameters: d=500, T=3, p=0.99

```

--- Series Logistic Regression ---
[0.0594 0.0584 0.0612 0.06 ]

```

```

--- Series Log-density Estimation ---
[0.0578 0.0592 0.0594 0.0608]

```

```

--- Simple Logistic Regression ---
[0.207 0.1464 0.1186 0.103 ]

```

## 2.2 Gaussian Mixtures in Low Dimensional Setting (d=2)

```
[95]: d = 2
p = 0.99
cfg = st.SelfTrainingConfig(
    N=N, M=M, N_test=N_test, d=d,
    mu_plus= 1 * np.ones(d), mu_minus= -1 * np.ones(d),
    p=p, T=T, n_runs=1, n_jobs=1, seed=seed,
    sigma_scale=sigma_scale
)

rng = np.random.default_rng(seed)
X, Y = st.sample_data(total, cfg, rng)
X_test, Y_test = st.sample_data(N_test, cfg, rng)
X_LO = X[:N]
Y_LO = Y[:N]
X_U0 = X[N:total]

results(X_LO, Y_LO, X_U0, X_test, Y_test, pi, p, T, K_seq)


# Data Visualization
plt.scatter(X[Y==1][:,0], X[Y==1][:,1], s=1)
plt.scatter(X[Y==0][:,0], X[Y==0][:,1], s=1)
plt.show()
```

Parameters: d=2, T=3, p=0.99

--- Series Logistic Regression ---  
[0.1588 0.1578 0.1576 0.159 ]

--- Series Log-density Estimation ---  
[0.159 0.1598 0.1654 0.1704]

--- Simple Logistic Regression ---  
[0.1676 0.1698 0.1706 0.1702]



## 2.3 Moons

```
[96]: p = 0.9
X, Y = make_moons(n_samples=total, noise=np.sqrt(sigma_scale) * 0.1,
                  random_state=seed)
X_test, Y_test = make_moons(n_samples=N_test, noise=np.sqrt(sigma_scale) * 0.1,
                             random_state=seed)

Y = 2 * Y - 1
Y_test = 2 * Y_test - 1
X_L0 = X[:cfg.N]
Y_L0 = Y[:cfg.N]
X_U0 = X[cfg.N:cfg.N + cfg.M]

results(X_L0, Y_L0, X_U0, X_test, Y_test, pi, p, T, K_seq)


# Data Visualization
plt.scatter(X[Y==-1][:,0], X[Y==-1][:,1], s=1)
plt.scatter(X[Y==1][:,0], X[Y==1][:,1], s=1)
plt.show()
```

Parameters: d=2, T=3, p=0.9

```
--- Series Logistic Regression ---  
[0.4264 0.2982 0.3012 0.3406]
```

```
--- Series Log-density Estimation ---  
[0.463  0.311  0.3104 0.279 ]
```

```
--- Simple Logistic Regression ---  
[0.1266 0.127  0.1272 0.126 ]
```



## 2.4 Circles

```
[97]: p = 0.9  
X, Y = make_circles(n_samples=total, noise=np.sqrt(cfg.sigma_scale) * 0.1,  
                    random_state=cfg.seed, factor=0.5)  
X_test, Y_test = make_circles(n_samples=cfg.N_test, noise=np.sqrt(cfg.  
    ↪sigma_scale) * 0.1,  
                              random_state=cfg.seed, factor=0.5)  
  
Y = 2 * Y - 1  
Y_test = 2 * Y_test - 1  
X_LO = X[:cfg.N]  
Y_LO = Y[:cfg.N]  
X_UO = X[cfg.N:cfg.N + cfg.M]
```

```

results(X_L0, Y_L0, X_U0, X_test, Y_test, pi, p, T, K_seq)

# Data Visualization
plt.scatter(X[Y==1][:,0], X[Y==1][:,1], s=1)
plt.scatter(X[Y==1][:,0], X[Y==1][:,1], s=1)
plt.show()

```

Parameters: d=2, T=3, p=0.9

```

--- Series Logistic Regression ---
[0.0482 0.049  0.0486 0.047 ]

```

```


--- Series Log-density Estimation ---
[0.0636 0.0492 0.054  0.056 ]

```

```

--- Simple Logistic Regression ---
[0.5286 0.5286 0.5286 0.5286]

```



[ ]: